



DEVOXXTM
the java™ community conference

G1 Garbage Collector



Jaromir Hamala <jhamala@c2b2.co.uk>
Senior Consultant
C2B2
@c2b2consulting



Jaromír Hamala

Who Am I?

- A curious developer
- Who likes to know more
 - About computers
 - World
 - And beer!



Agenda

- Memory Management
- Garbage Collectors in JVM
- CMS vs. G1
- G1 Tuning
- Recommendations

Manual Memory Management I

- Not all the memory is yours!
- You need to ask for it

malloc() -> sbrk()

- And return when you no longer need it

free() -> sbrk()

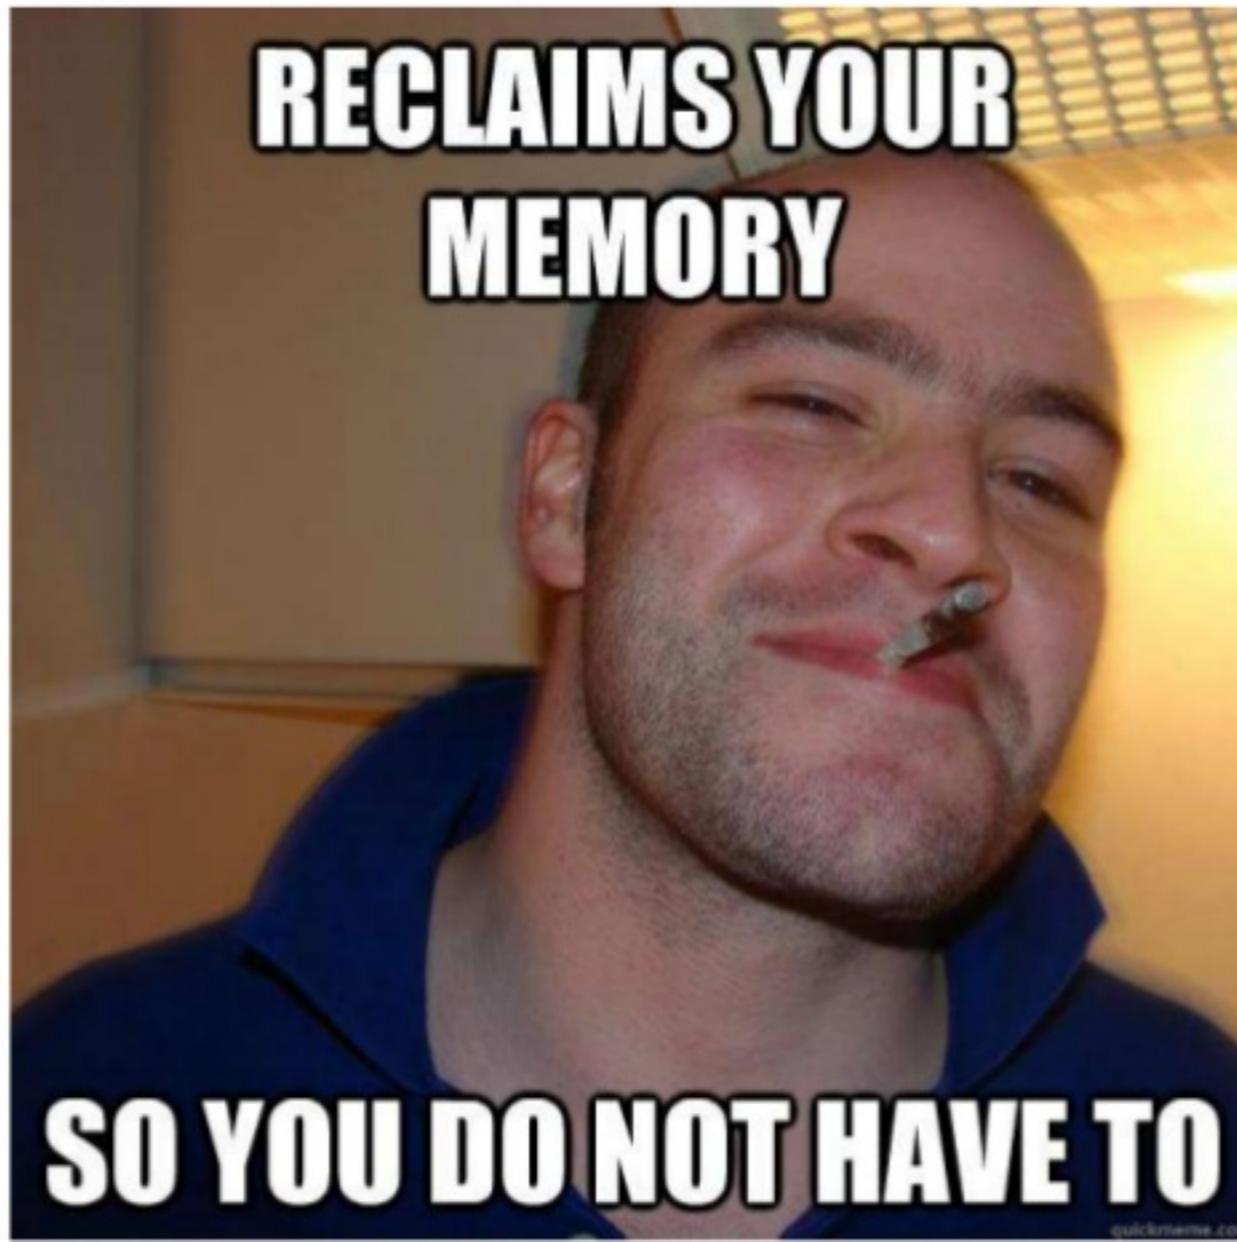


Manual Memory Management II

- Asking for something is easy
- Returning is way more complicated!

(we all know it!)

Garbage Collector I



A little bit of history

- **3000 B.C.** – “*The first landfill is developed when Knossos, Crete digs large holes for refuse. Garbage is dumped and filled with dirt at various levels.*”
- In LISP from end of '50s!
- In Java from the 1st version



Requirements on GC

- High Throughput
- Low Latency
- Minimal Footprint
- Reliable
- Easy to use/tune

Generational Hypothesis

- Most objects die young
- Old objects rarely reference young objects

Generational Garbage Collector

- Split the heap into regions
 - Create new objects in Young
 - Move mature objects to Old
-
- Different strategies for different regions

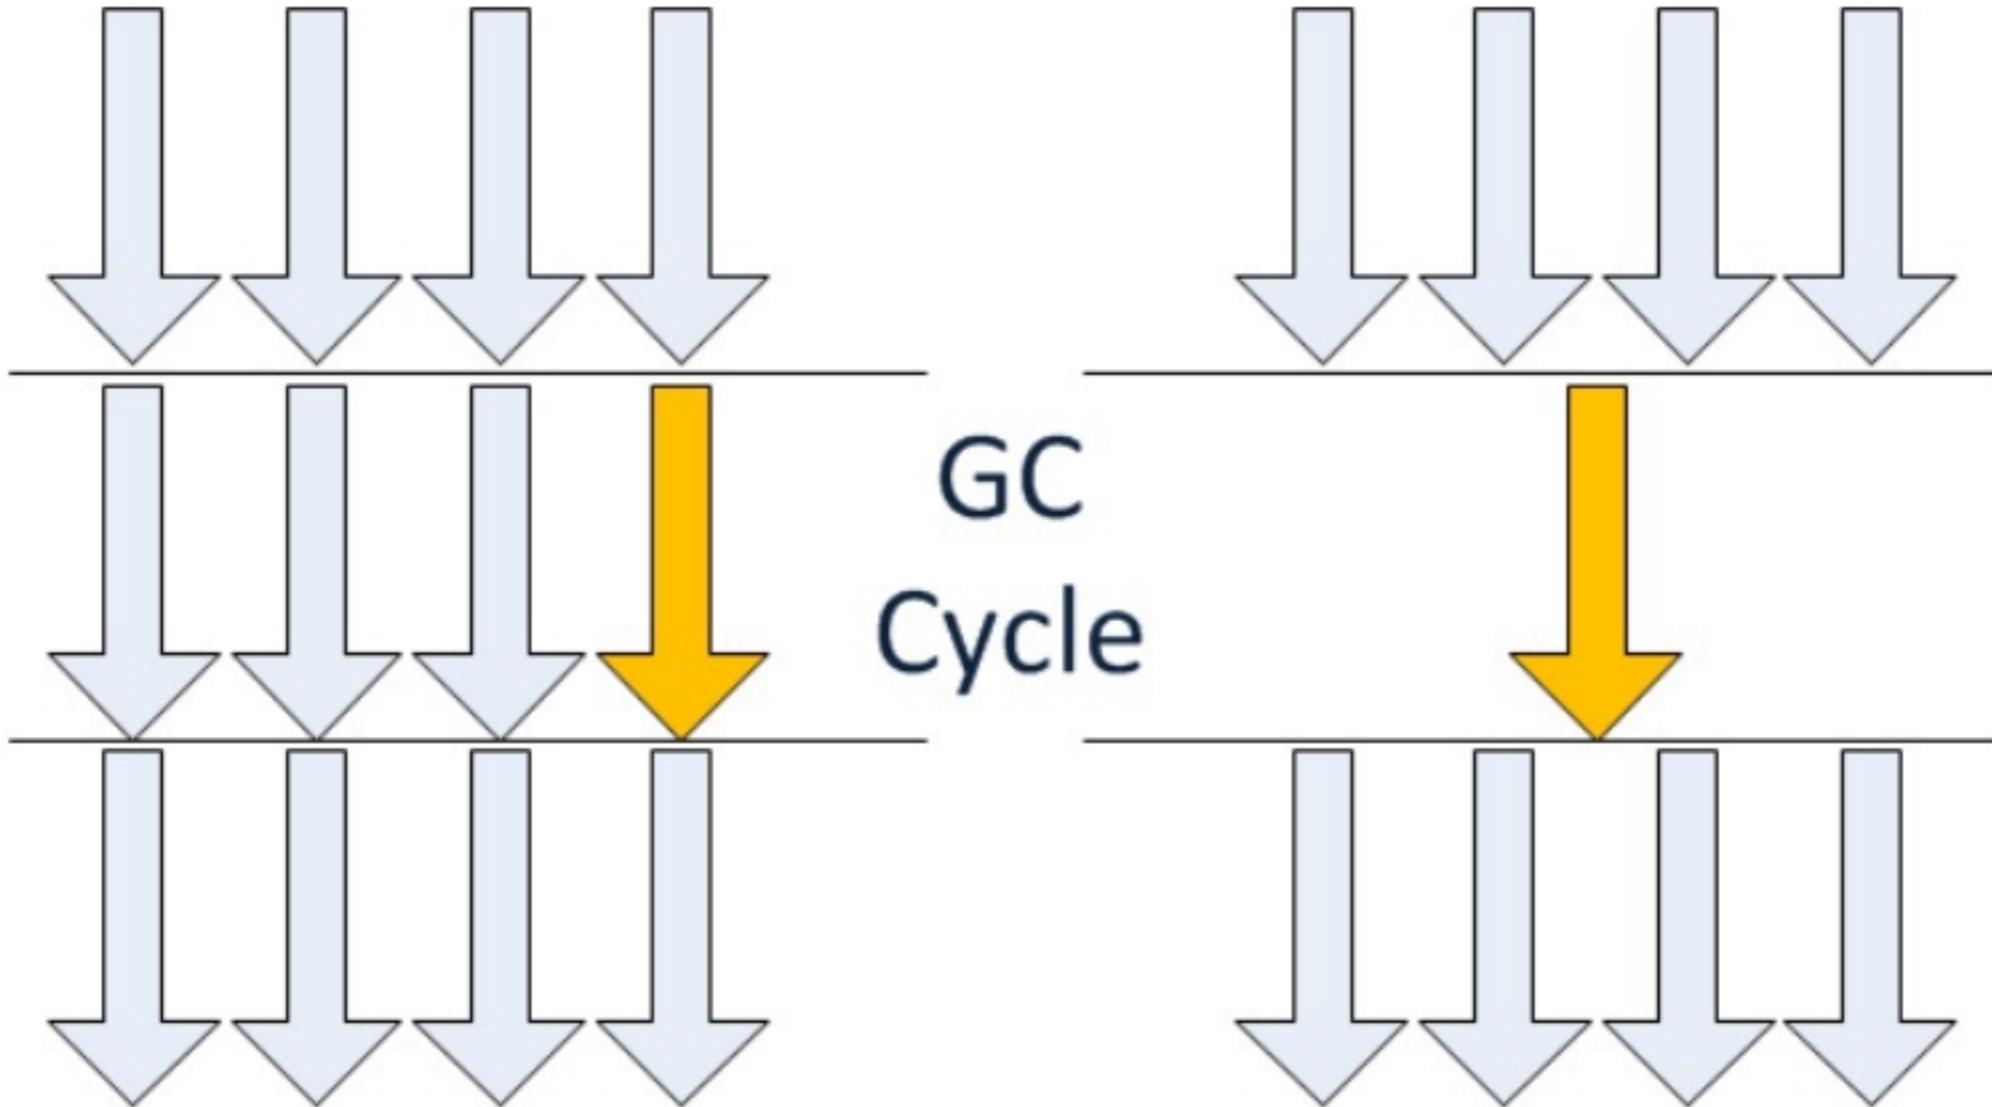
Young

Old

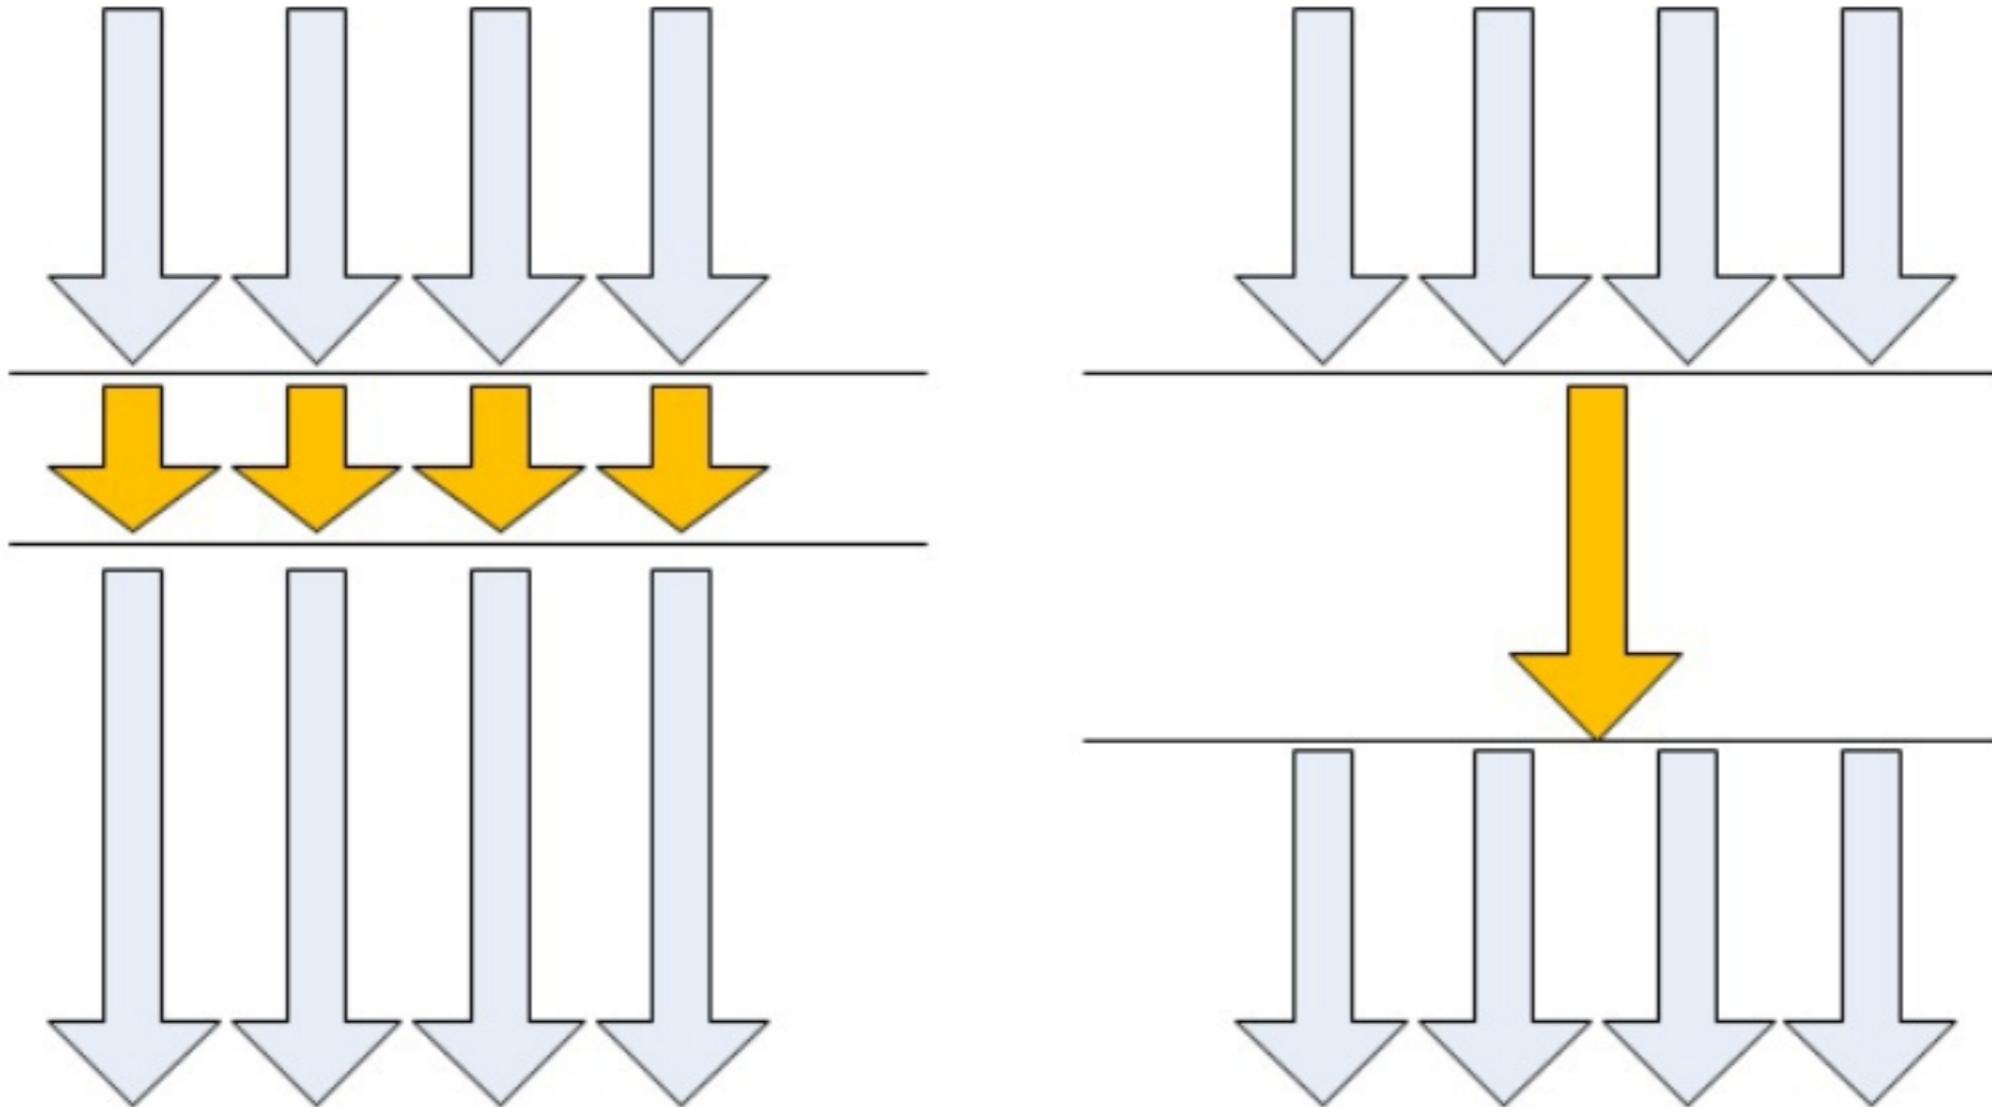
Characteristics of GCs

- Concurrent
- Parallel
- Compacting

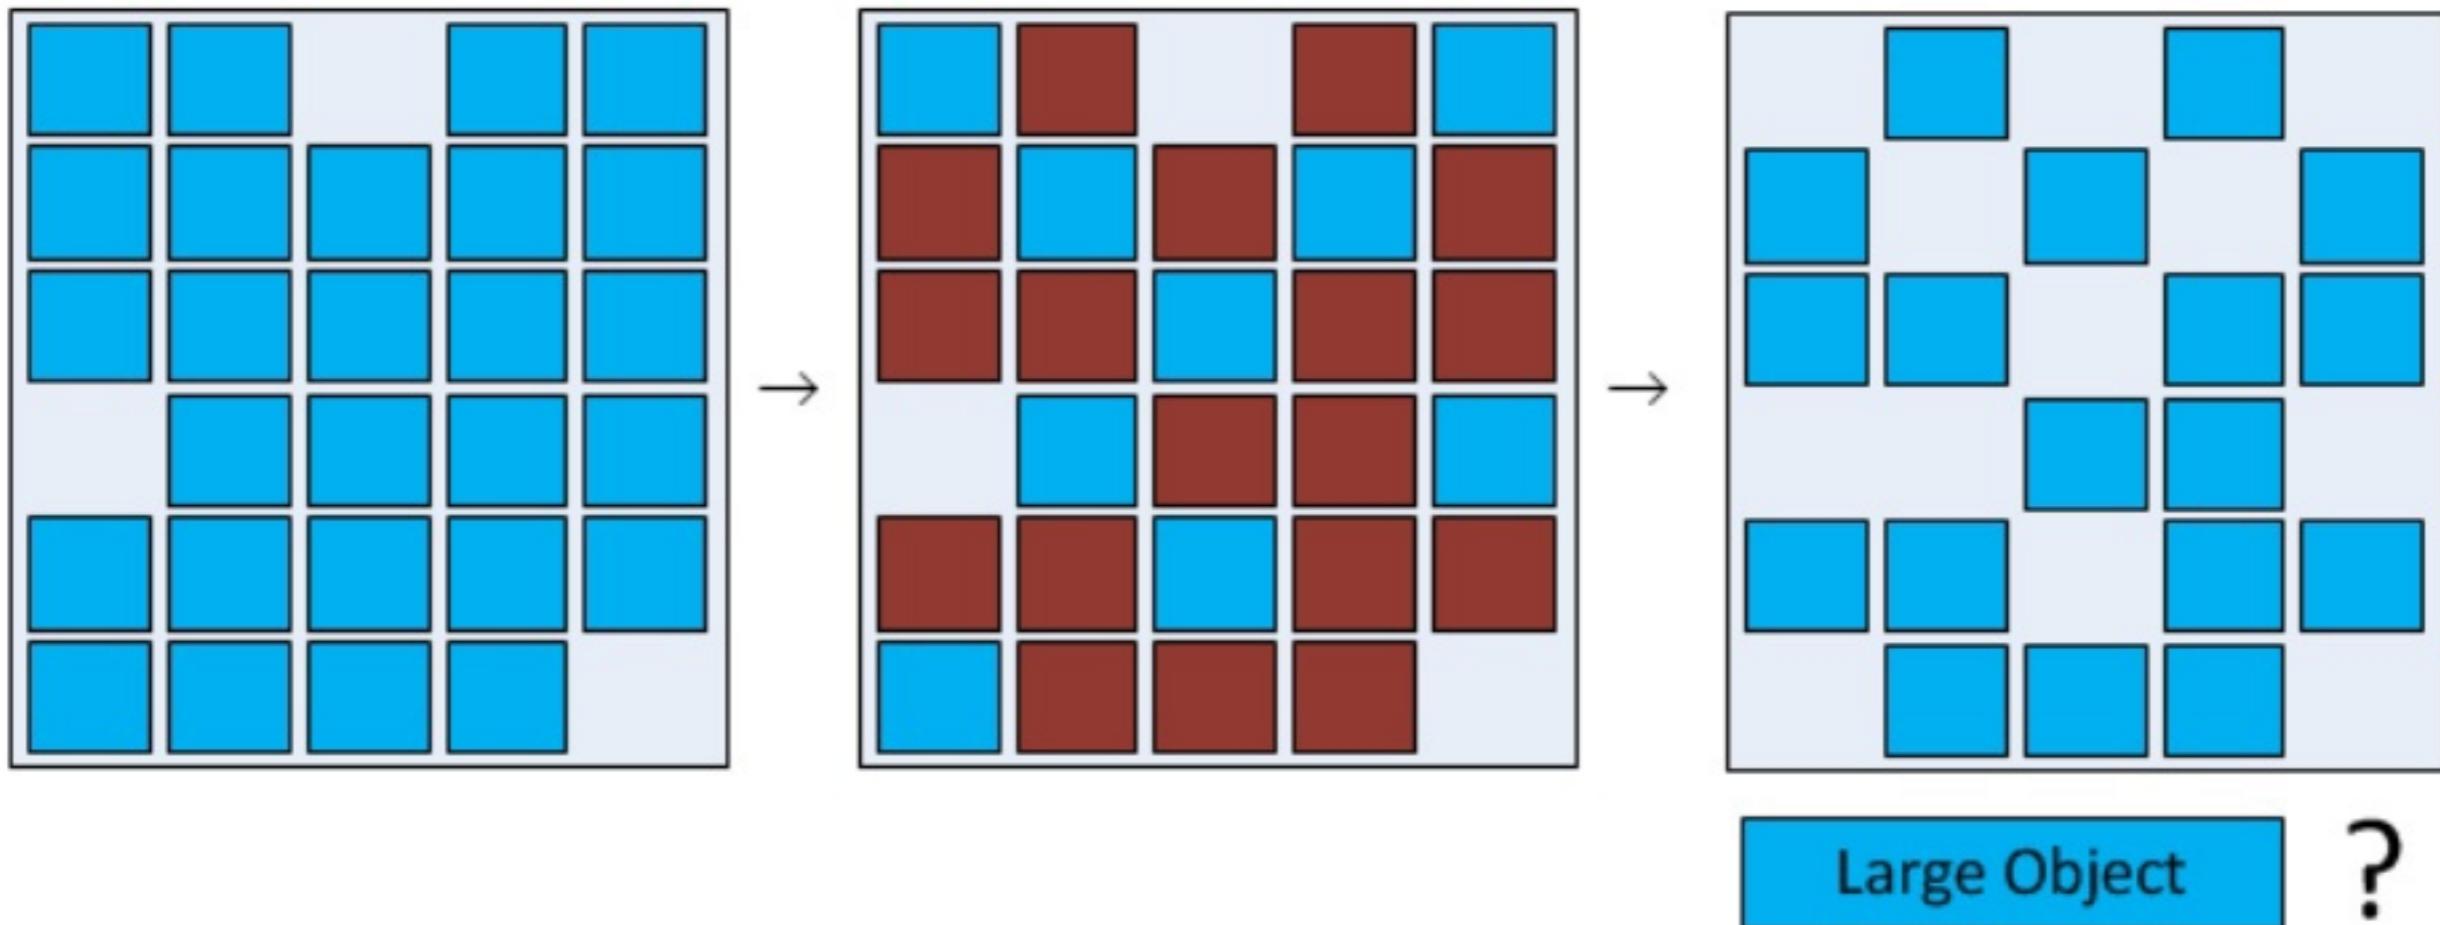
Concurrent GC



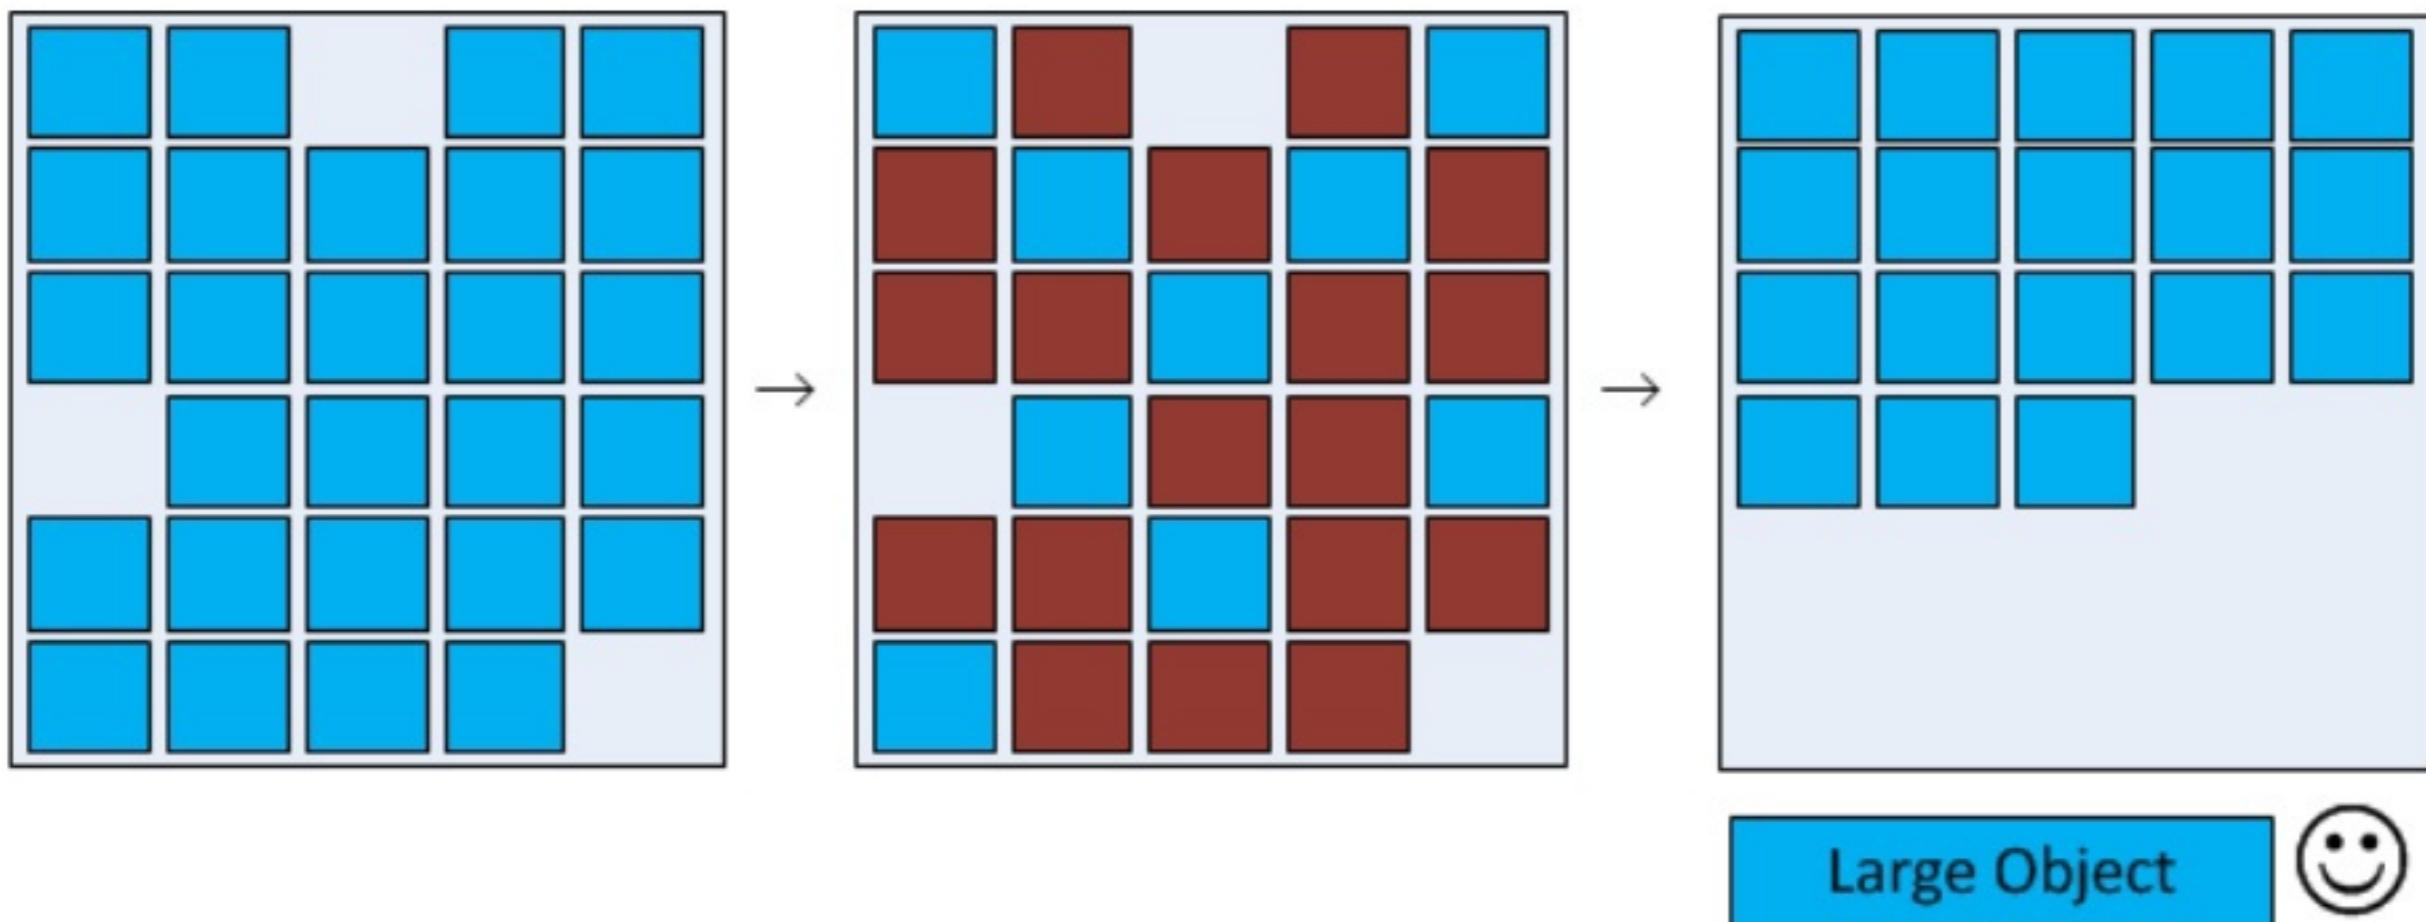
Parallel GC



Heap Fragmentation

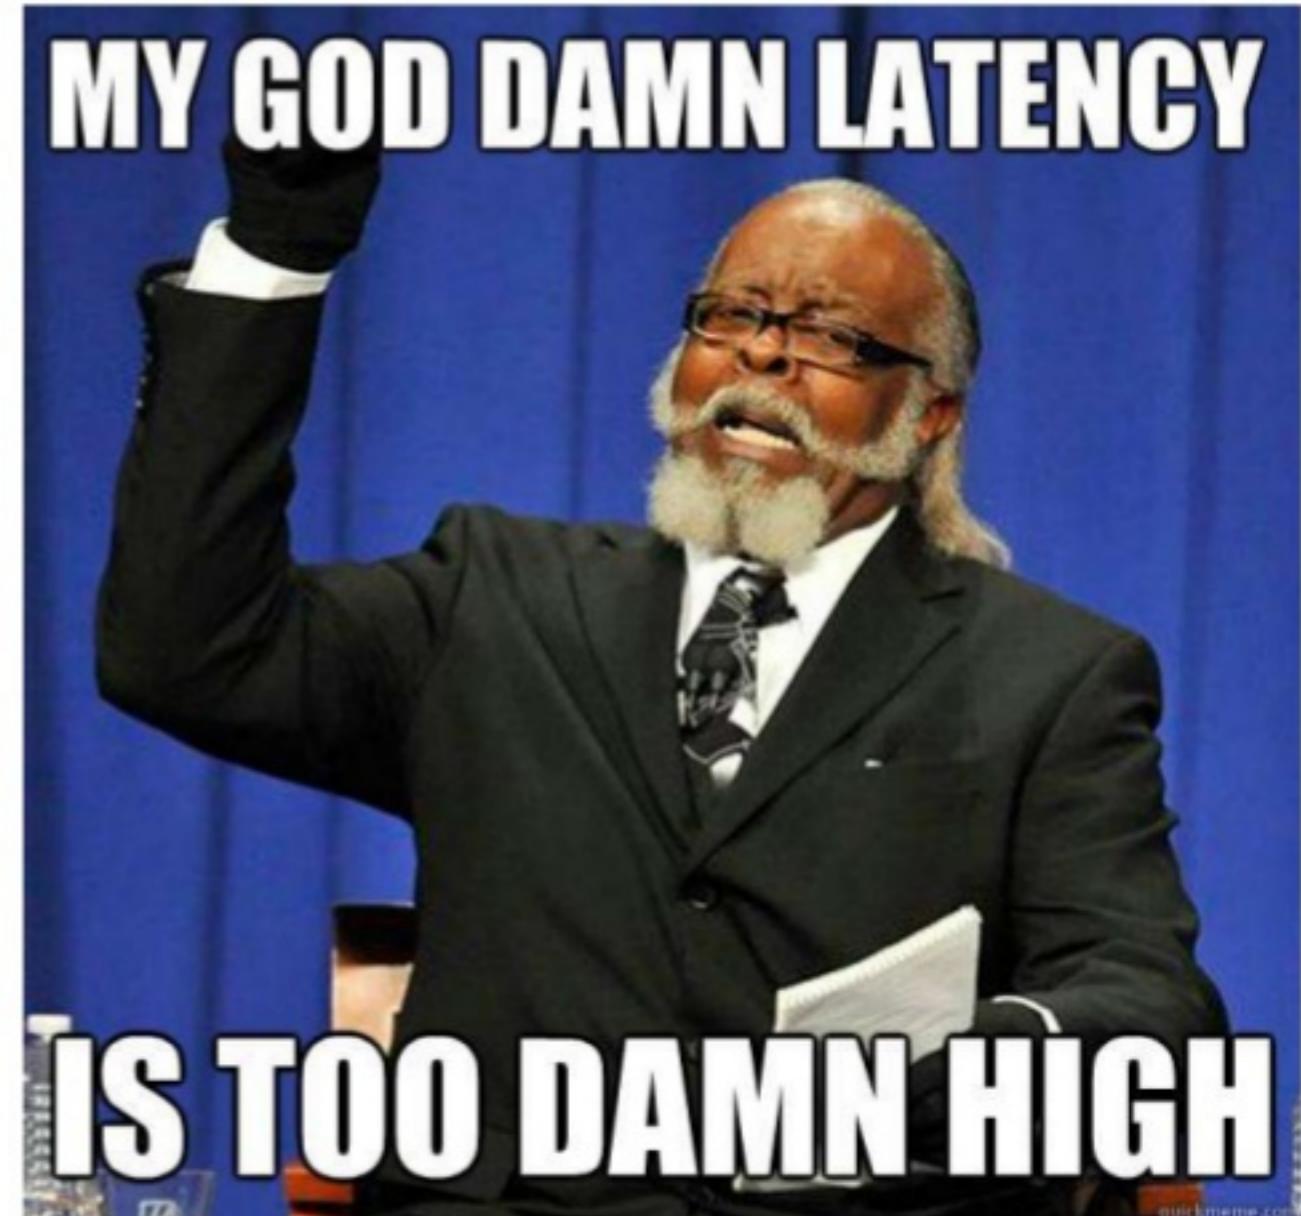


Heap Compacted



What we have in JVM? I

- Serial GC
- Parallel GC

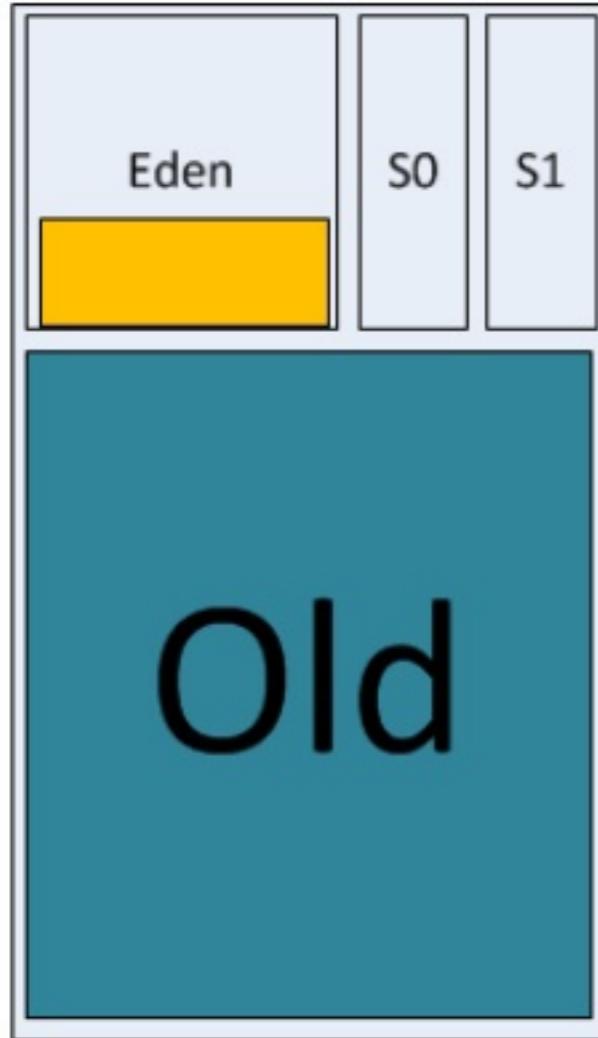


What we have in JVM? II

- CMS
- G1
 - since 6u14
 - supported since 7u4



How CMS works in young generation II

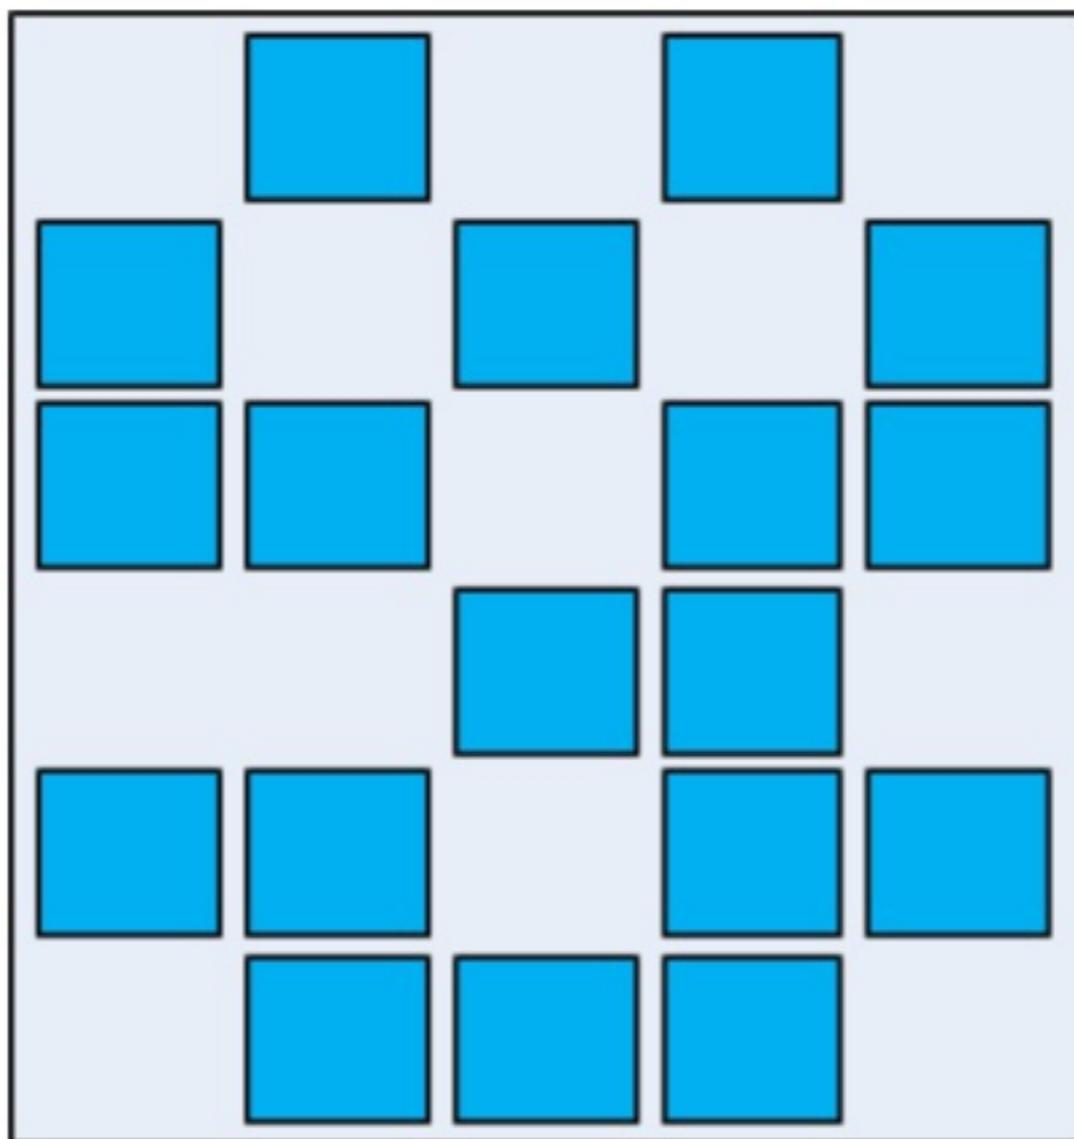


How CMS works in young generation II

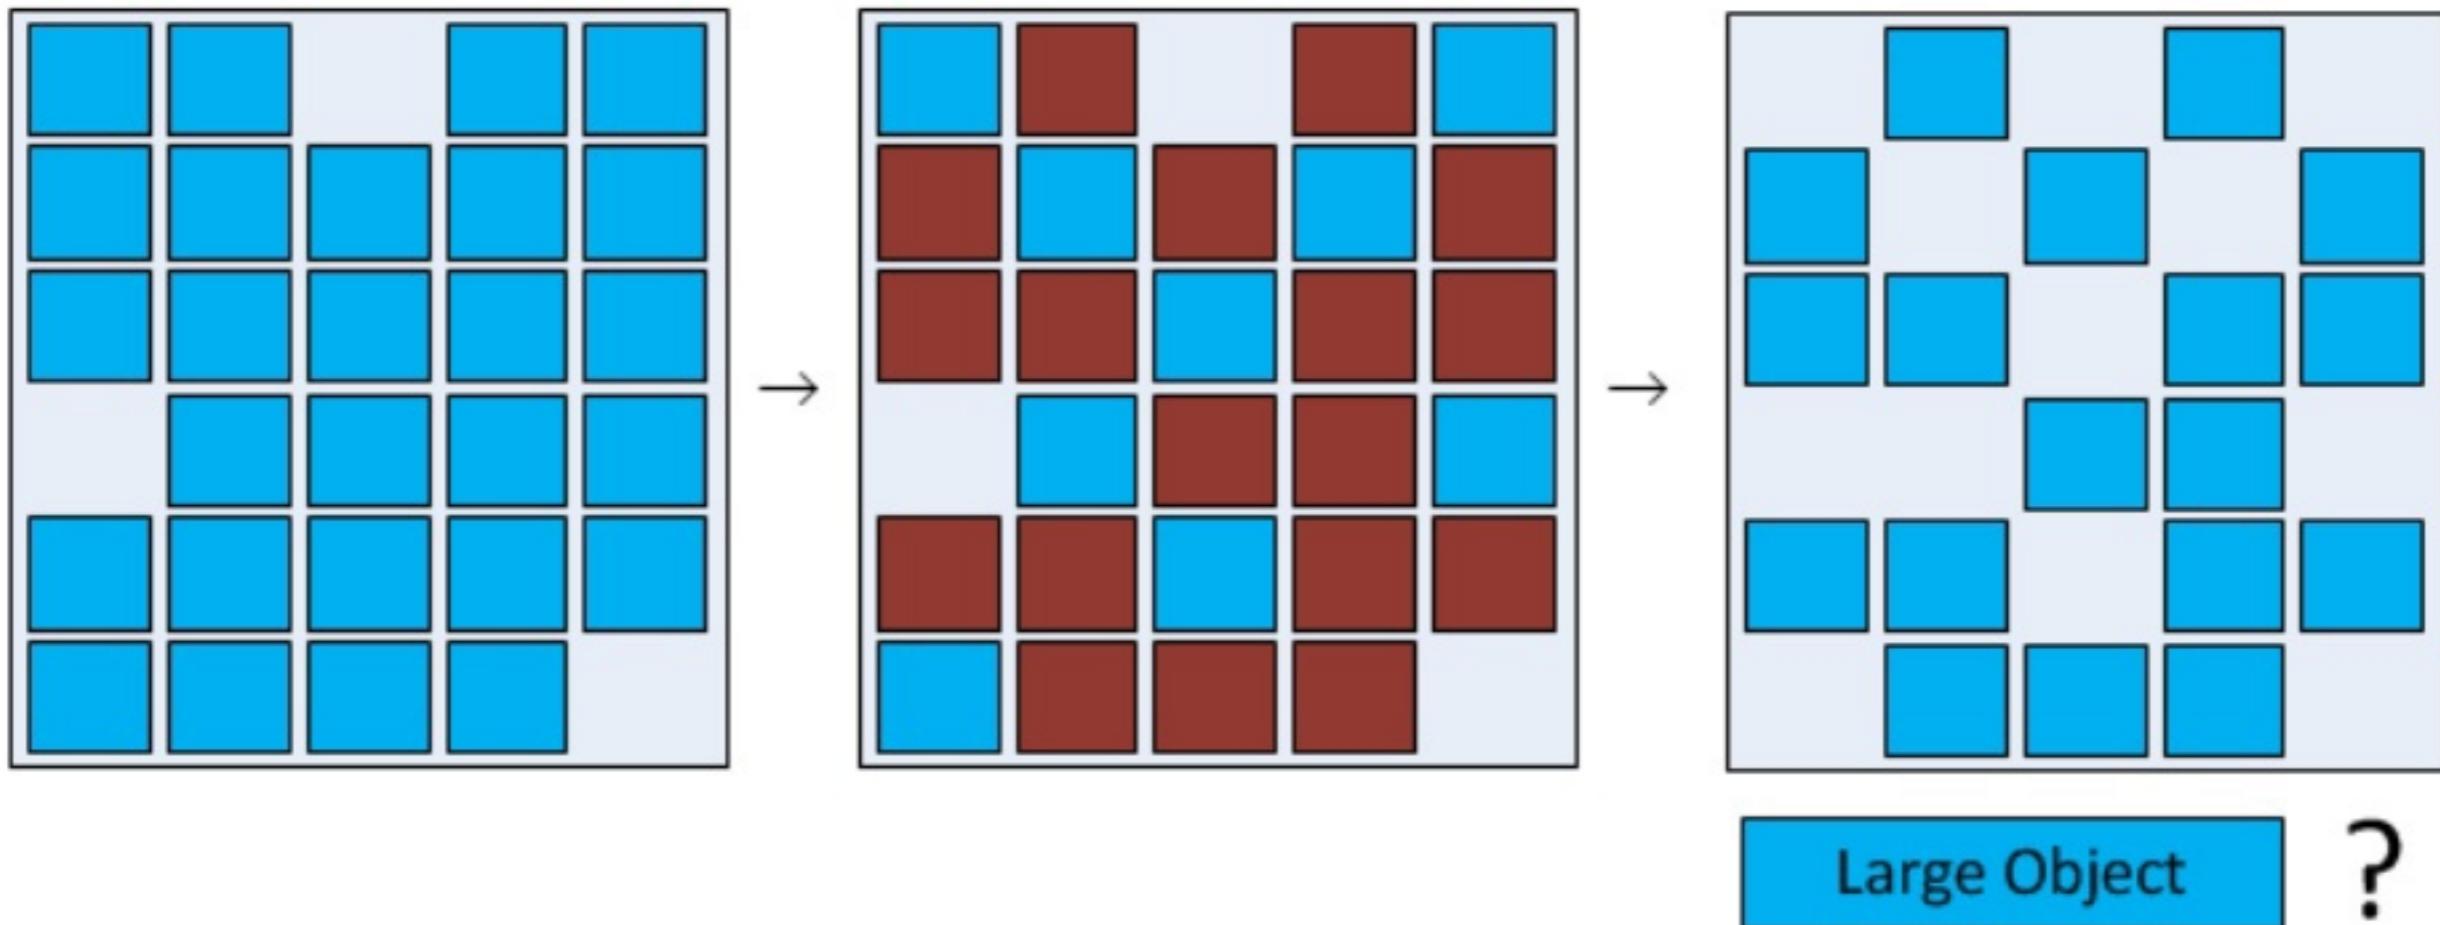


How CMS works in old generation

- Main phases of CMS:
 - Initial Mark (STW)
 - Concurrent Mark
 - Remark (STW)
 - Concurrent Sweep
- Features:
 - Continuous block of memory
 - No compacting!



Heap Fragmentation



What's the catch?

- Fragmentation
- Remark phase (STW)
- Very Large Heaps

“It’s working very well except when it’s not.” –M. Brasier

Promotion Failure
Concurrent Mode Failure



Stop The World



G1

- Concurrent
- Parallel
- Compacting

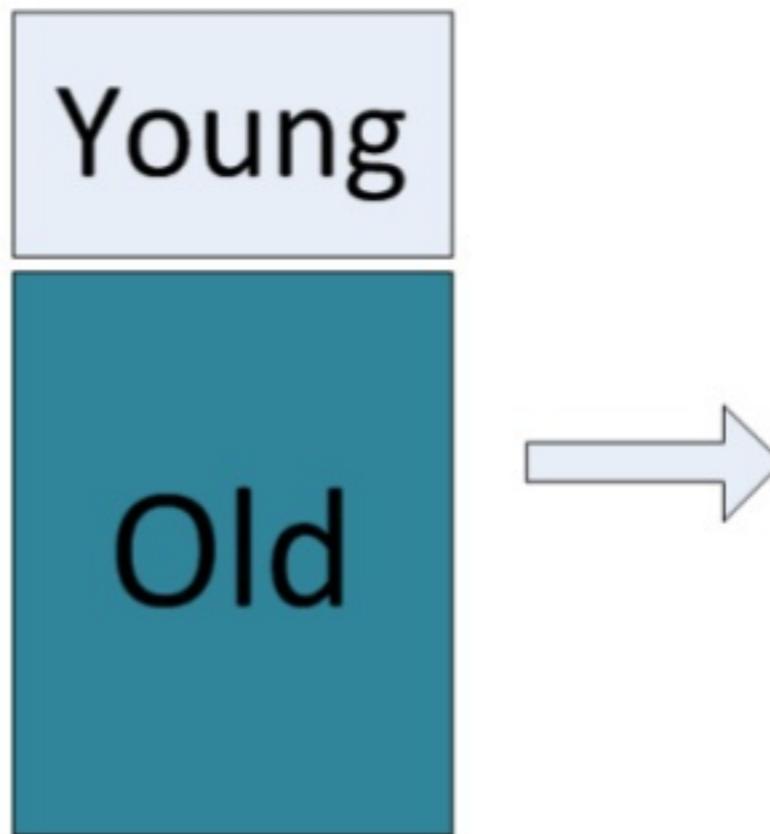


G1 Goals

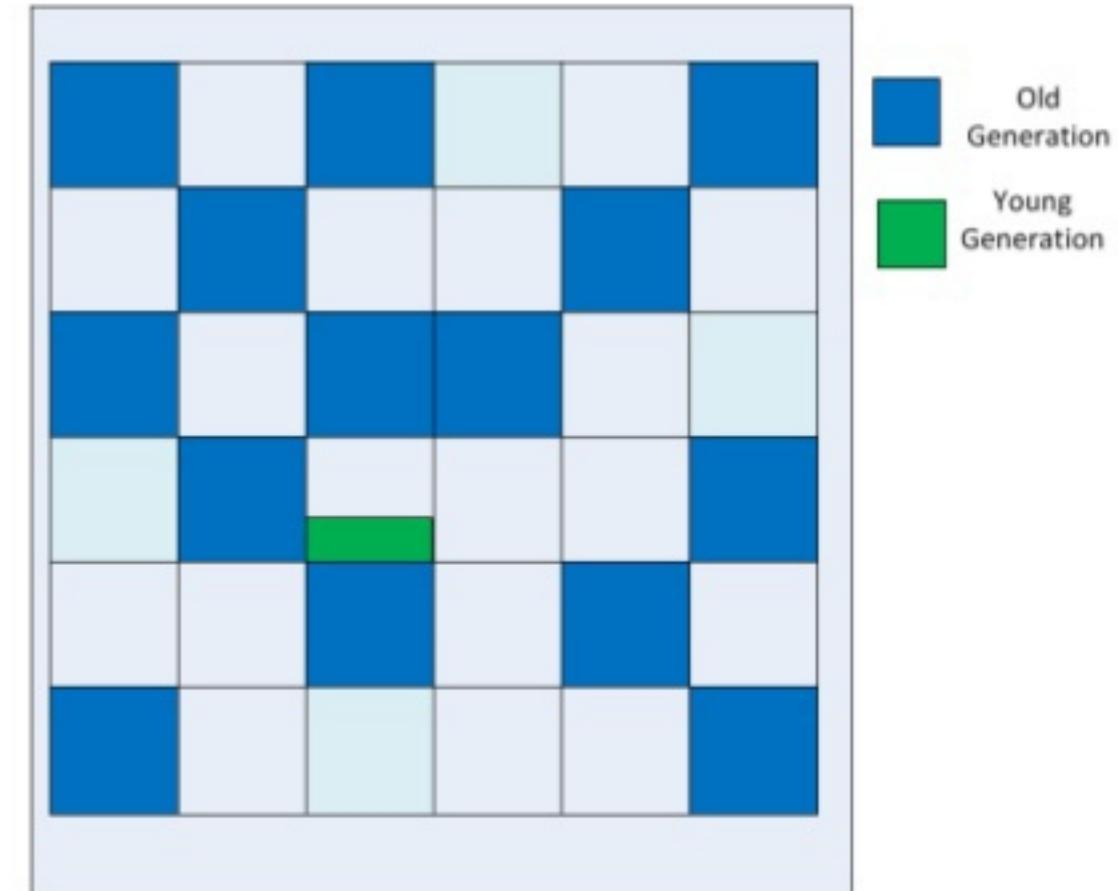
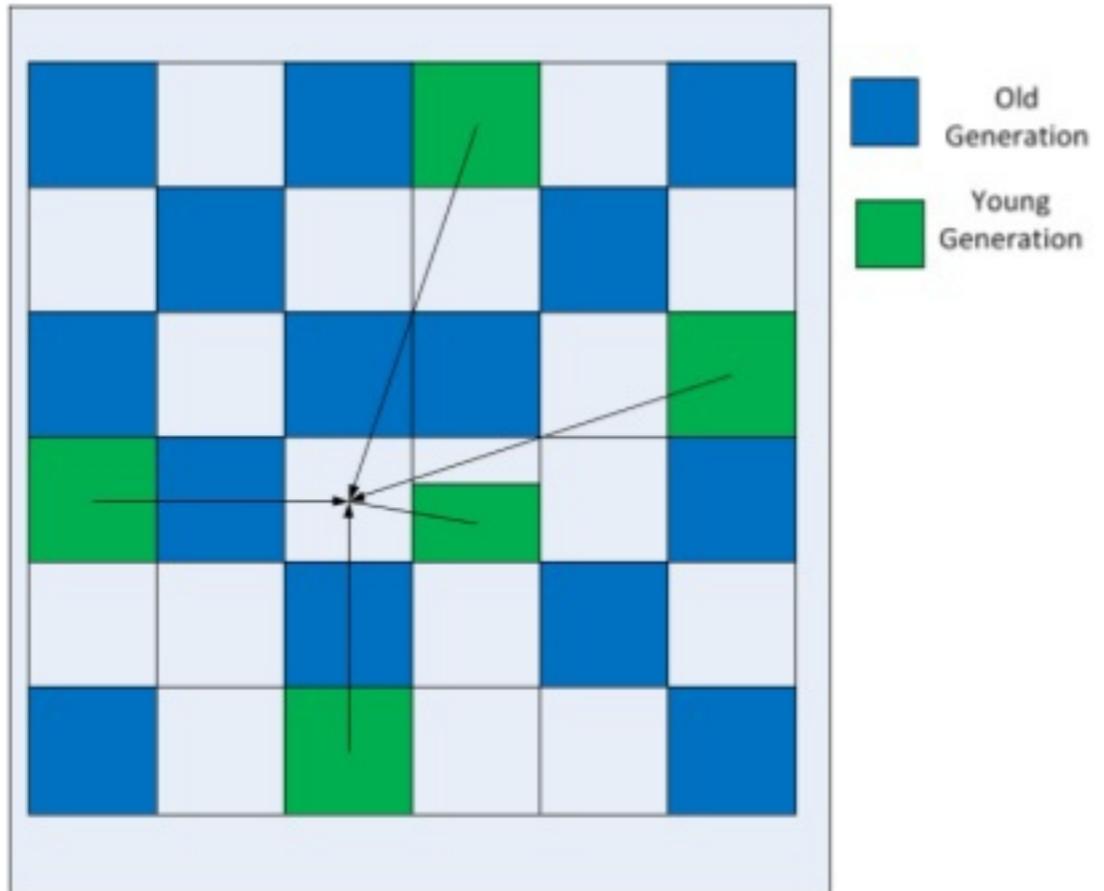
- Low Latency
- Better Predictability
- Easy to use & tune

G1 Heap Layout

- Heap divided into +- 2000 regions of equal size
- No physical separation between young / old generation
- Each region can be either young or old
- Object moved between regions during collections
- Humongous Region
 - For large objects
 - Expensive collecting as of now!



G1 Young Generation GC



G1 Old Generation GC I.

- *Combination of CMS and Parallel Compacting GC*

- Concurrent Marking Phase

- Remark still STW

1929.299: [GC remark 1929.299: [GC ref-proc, 0.0048180 secs], 0.0258670 secs]
[Times: user=0.03 sys=0.00, real=0.02 secs]

- Initial Marking part of evacuation pause

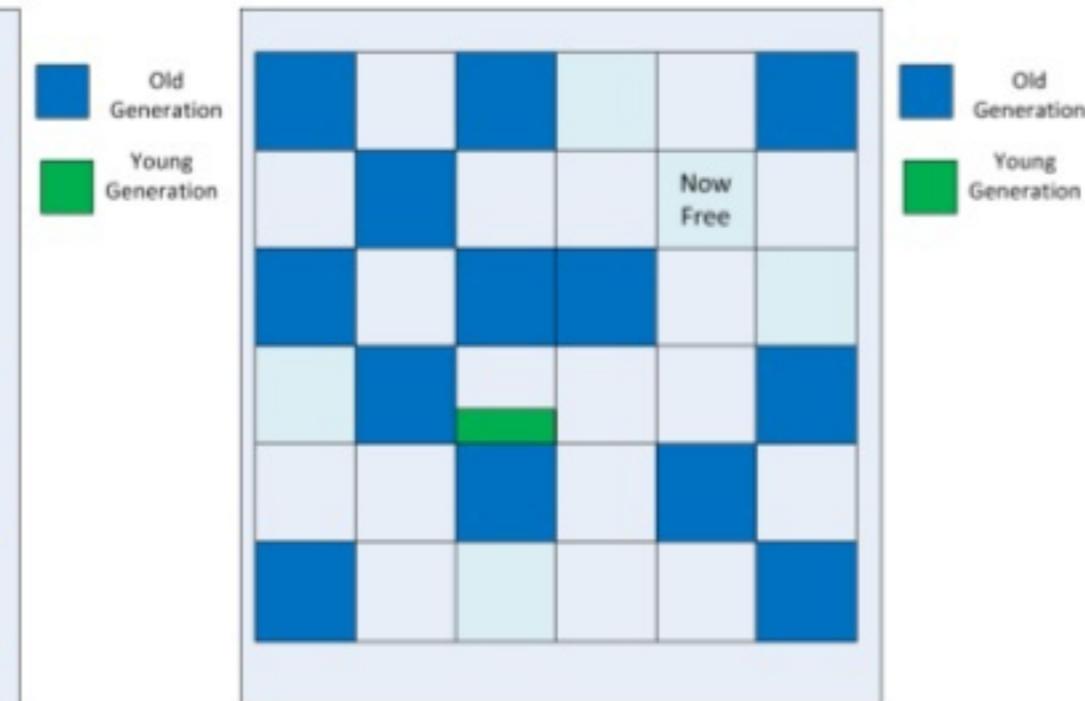
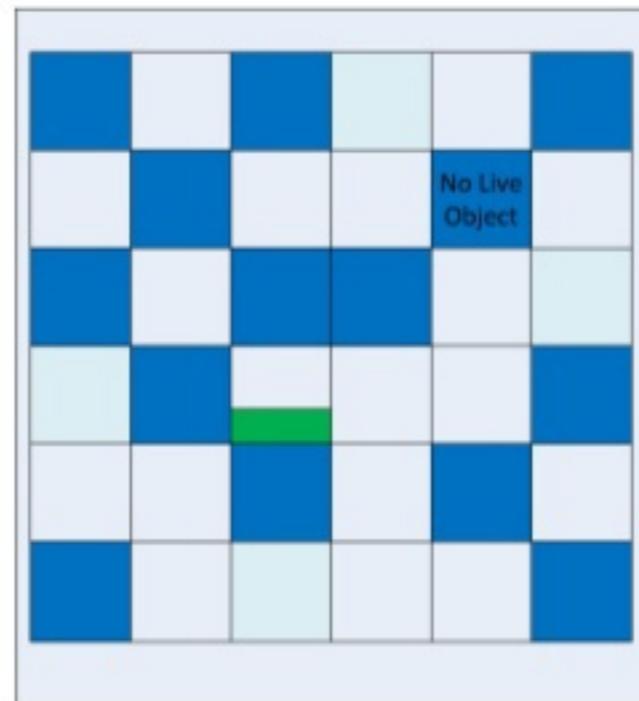
62.583: [GC pause (young) (initial-mark), 0.09812400 secs]

- Calculating statistics data

- Ratio of living objects inside region
- Who references our region

G1 Old Generation GC II.

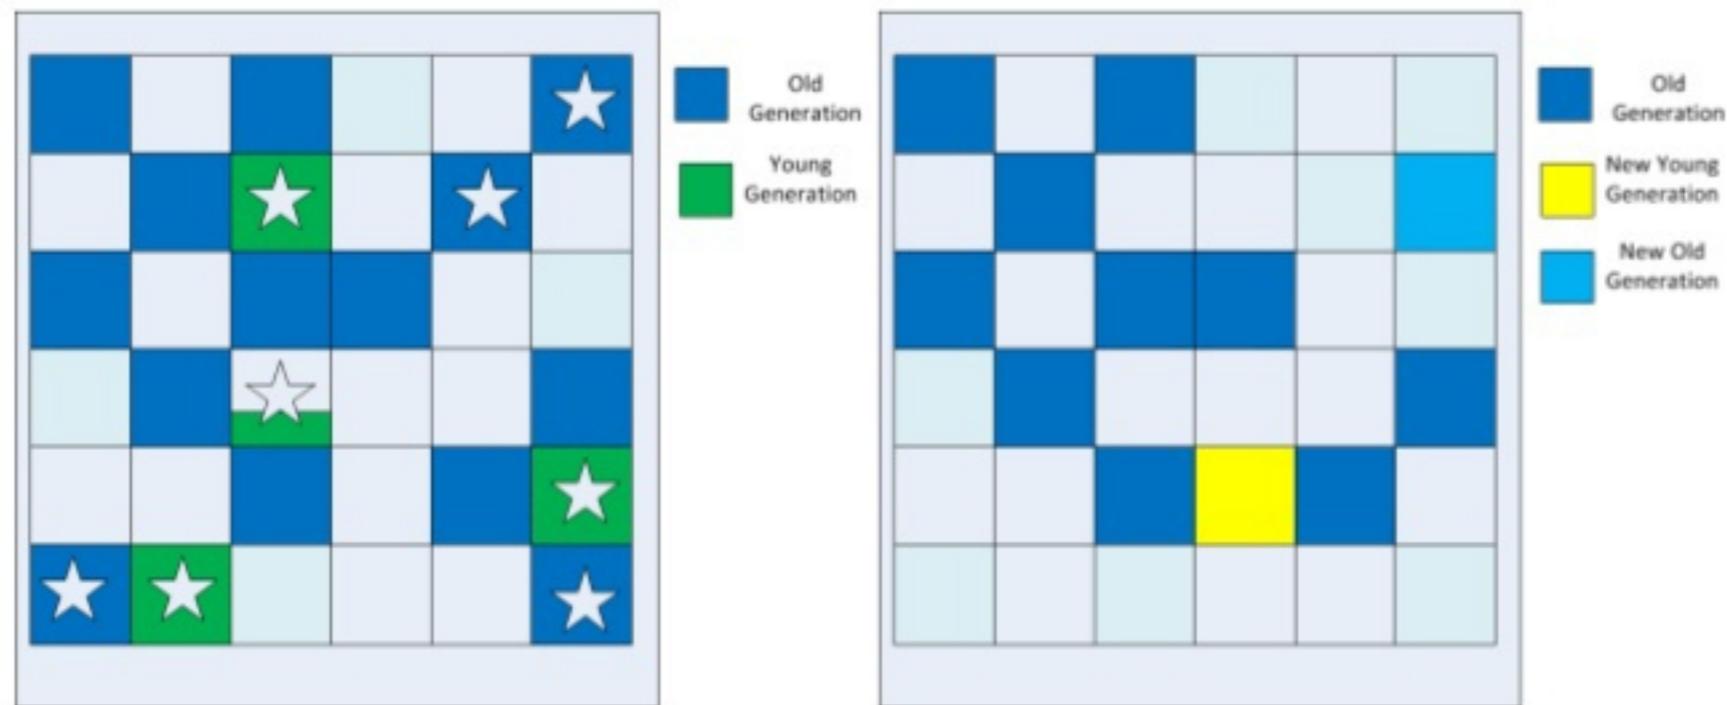
- Regions with no living objects can be reclaimed immediately (after remark)



G1 Old Generation GC III.

- Regions with smallest ratio of living objects chosen to evacuate during next GC

53.695: [GC pause (mixed), 0.05568800 secs]



G1 Old Generation GC IV.

- Concurrent Marking Phase
 - Calculate liveness ratio
 - Find *best* regions to evacuate

- Reclaiming Regions
 - After Remark
 - During evacuation pause

Remember Set

- Structure to track references from other regions
- One per region
- Allows to collects objects without tracking whole heap to find references
- Maintenance based on write barrier

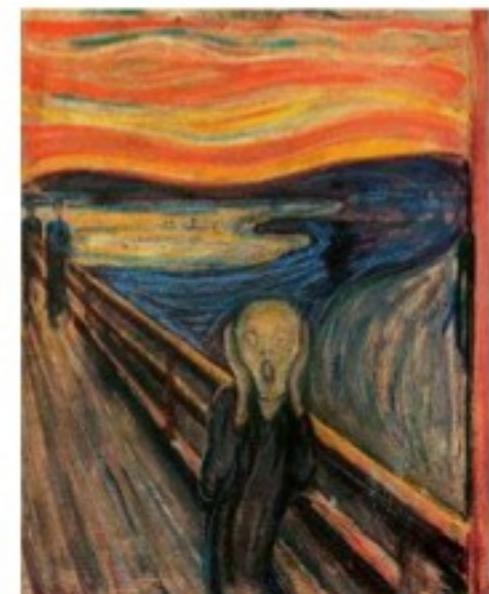
There is no free lunch!

Gotcha! (Part I.)

- More inter-region references -> Bigger Remember Set
- Large Remember Set -> Slow G1
- Remember: Large Object -> Humongous Region -> Collecting not optimized as yet

G1 Tuning Fundamentals

- Main goal of G1 is a low latency
 - XX:+UseCMSInitiatingOccupancyOnly
- If you can tolerate latency use Parallel GC
 - XX:CMSInitiatingOccupancyFraction=60
- Another goal is simplicity (of tuning)
 - XX:CMSTriggerRatio=60
 - XX:CMSWaitDuration=1500
- Most important tuning option:
 - XX:MaxGCPauseMillis=300
 - XX:ConcGCThreads=2
 - XX:SurvivorRatio=8
 - Xmn1g
 -
- Influence maximum amount of work per collection
- No hard promise
- Best effort only



Other Import Tuning Options I

-XX:InitiatingHeapOccupancyPercent=n

- When to start a concurrent GC cycle
- Percent of an entire heap, not just old generation!

Too Low -> Unnecessary GC overhead

Too High -> “Space Overflow” -> Full GC

Other Import Tuning Options II

-XX:G1OldCSetRegionLiveThresholdPercent=n

- Threshold for region to be included in a Collection Set
- To High -> More aggressive collecting
 - More live objects to copy
- To Low -> Wasting some heap

Other Import Tuning Options III

-XX:G1MixedGCCountTarget=n

- How many Mixed GC / Concurrent Cycle
- To High -> Unnecessary overhead
- To Low -> Longer pauses

Considerations Before Tuning

- Some options cause your PauseTimeTarget to be ignored!
 -Xmn
- Fixed young generation size

Logging the GC Activity

- Verbose GC is your friend!
`-XX:+PrintGC -XX:+PrintGCDetails -
XX:+PrintGCTimeStamps`
- Always use it! Even in production.
- Very low overhead. There is no excuse!
- Will help you analyze help your GC is performing

G1 GC Log

```
1.297: [GC pause (young), 0.02828800 secs]
[Parallel Time: 24.2 ms]
[GC Worker Start (ms): 1296.7 1302.7
 Avg: 1299.7, Min: 1296.7, Max: 1302.7, Diff: 6.0]
[Ext Root Scanning (ms): 16.5 2.9
 Avg: 9.7, Min: 2.9, Max: 16.5, Diff: 13.6]
[Update RS (ms): 0.0 0.0
 Avg: 0.0, Min: 0.0, Max: 0.0, Diff: 0.0]
[Processed Buffers : 0 9
 Sum: 9, Avg: 4, Min: 0, Max: 9, Diff: 9]
[...]
```

G1 GC Log

1.297: [GC pause (young), 0.02828800 secs]

- Young GC
- Started 1.297s from the start of JVM
- Total elapsed time was 0.02828800s

[Parallel Time: 24.2 ms]

- Total elapsed time spent by parallel GC thread

G1 GC Log

```
[GC Worker Start (ms): 1296.7 1302.7  
          Avg: 1299.7, Min: 1296.7, Max: 1302.7, Diff: 6.0]
```

- Two threads and starting times

```
[Ext Root Scanning (ms): 16.5 2.9  
          Avg: 9.7, Min: 2.9, Max: 16.5, Diff: 13.6]
```

- Time spent by each thread by scanning the roots

G1 GC Log

```
[Update RS (ms):  0.0  0.0
    Avg:  0.0, Min:  0.0, Max:  0.0, Diff:  0.0]
    [Processed Buffers : 0 9
        Sum: 9, Avg: 4, Min: 0, Max: 9, Diff: 9]
```

- Time spent by updating Remember Sets regions
- Update Buffers is a structure tracking changes in references

```
[Scan RS (ms):  0.0  0.1
    Avg:  0.0, Min:  0.0, Max:  0.1, Diff:  0.1].
```

- Time spent by scanning areas from Remember Set

G1 GC Log

[Object Copy (ms): 7.6 15.1
Avg: 11.3, Min: 7.6, Max: 15.1, Diff: 7.5]

- Time spent by copying objects to other regions
- Usually significant portion of a total time

[Termination (ms): 0.0 0.0
Avg: 0.0, Min: 0.0, Max: 0.0, Diff: 0.0]
[Termination Attempts : 1 1
Sum: 2, Avg: 1, Min: 1, Max: 1, Diff: 0]

- Time spent by offering to terminate
- Threads are stealing work from the others

G1 Logging Options I

- -XX:+PrintAdaptiveSizePolicy

53.628: [GC pause (mixed) 53.628: [G1Ergonomics (CSet Construction) start choosing CSet, predicted base time: 54.79 ms, remaining time: 145.21 ms, target pause time: 200.00 ms]

53.628: [G1Ergonomics (CSet Construction) add young regions to CSet, eden: 51 regions, survivors: 2 regions, predicted young region time: 9.75 ms]

53.628: [G1Ergonomics (CSet Construction) finish adding old regions to CSet, reason: old CSet region num reached max, old: 27 regions, max: 27 regions]

53.628: [G1Ergonomics (CSet Construction) finish choosing CSet, eden: 51 regions, survivors: 2 regions, old: 27 regions, predicted pause time: 131.56 ms, target pause time: 200.00 ms]

53.700: [G1Ergonomics (Concurrent Cycles) do not request concurrent cycle initiation, reason: still doing mixed collections, occupancy: 127926272 bytes, allocation request: 0 bytes, threshold: 125986365 bytes (45.00 %), source: end of GC]

G1 Evacuation Failure

- Evacuation Failure Happens when we do not have empty regions
 - Heap is already at maximum size
 - Very expensive operation!
 - Search for “to-space overflow”
-
- Remedy:
 - Increase to heap size
 - -XX:InitiatingHeapOccupancyPercent=
 - More marking threads

Is G1 Production Ready?

- It depends!
 - Well, some people use it in production
- The days of G1 crashing JVM are hopefully over
- Still quite significant changes
- Bugs are still being found
 - *Bug 7181612 : G1: Premature Evacuation (7143858) strikes again*
 - *Bug 7143858 : G1: Back to back young GCs with the second GC having a minimally sized eden*
- You should probably evaluate/use G1 when everything else is failing
- And you have no alternative!

GC Decisioning: Rules of Thumb I

- Do I really need a Low Latency
 - No? Use Parallel GC
- Do I really need a Big Heap?
 - Yes? Think again!
- Do I really need a Big Heap?
 - No? Use small heap & Parallel GC
 - Yes? Try CMS 1st

GC Decisioning: Rules of Thumb II

- Is CMS performing well?
 - Yes? Done!
 - No? Tune it!
- Is tuned CMS performing well?
 - Yes? Done!
 - No? Tune it more!
- Is CMS performing well now?
 - Yes? Done!
 - No? Do you **really** need such big heap?

GC Decisioning: Rules of Thumb III

- Yes, I really need a Big Heap a Low Pauses!
 - This is the moment where you should start considering to use G1!
- In Nutshell:
 - For now treat G1 as a last resort
 - When other GCs are failing
 - Be aware of gotchas:
 - Very Large Objects
 - Large Remember Sets
 - Footprint?
 - Test carefully before using in a production!

G1 – Recent Commits

date: Tue Oct 30 11:45:51 2012 -0700

summary: 7200261: G1: Liveness counting inconsistencies during marking verification

date: Thu Sep 20 09:52:56 2012 -0700

summary: 7190666: G1: assert(_unused == 0) failed: Inconsistency in PLAB stats

date: Tue Aug 21 14:10:39 2012 -0700

summary: 7185699: G1: Prediction model discrepancies

G1 – Recent Commits – Conclusion?

- Do not use anything before 7u4
 - Use 7u9 or later
 - The latest version in the best case
- Do not use G1 from Java 6

What Are The Alternatives?

- Workaround / Hacks:
- Off Heap Allocations
 - DirectBuffers
 - JNI
- Object Pooling



Further Resources

- JavaOne 2012 G1 Talk, Charlie Hunt, Monica Beckwith
- hotstop-gc-use mailinglist
- <http://www.quora.com/What-is-the-status-of-the-implementation-of-the-G1-garbage-collector-for-the-JVM#>
-



Q&A